

Introduction to Optimization

Week 3 Assignment Solution

1a. Explain (simply) why $\sqrt{\pi}$ is the minimum of this function

$x^2 + y^2$ will always be greater than or equal to zero (positive) for all values of x and y .

Therefore:

$$x^2 + y^2 + \pi = (\text{at least}) \pi.$$

The function:

$$f(x, y) = x^2 + y^2 + \pi,$$

will have its minimum value when $x^2 + y^2$ is minimized, which happens when $x=0$ and $y=0$.

Hence, the minimum value of $f(x, y)$ is π , which occurs at the point $(0, 0)$.

So, π is the minimum value of the function $f(x, y)$.

1b. Compute the gradients (partial derivatives) of f with respect to x and y .

$$f(x) = \sqrt{x^2 + y^2 + \pi}$$

$$= \text{let } g = x^2 + y^2 + \pi$$

$$= f(x) = \sqrt{g}$$

With chain rule, f with respect to x and y becomes:

$$= \frac{df}{dx} = \frac{d\sqrt{g}}{dg} \cdot \frac{dg}{dx}$$

$$= \frac{df}{dy} = \frac{d\sqrt{g}}{dg} \cdot \frac{dg}{dy}$$

Differentiate \sqrt{g} with respect to g

$$= \frac{d\sqrt{g}}{dg} = \frac{1g^{\frac{1-2}{2}}}{2} = \frac{1g^{\frac{-1}{2}}}{2} = \frac{1}{2\sqrt{g}}$$

Partially differentiate $g = x^2 + y^2 + \pi$ with respect to x

$$= \frac{dg}{dx} = 2x$$

Partially differentiate $g = x^2 + y^2 + \pi$ with respect to y

$$\frac{dg}{dx} = 2y$$

Now differentiate $\frac{df}{dx} = \frac{d\sqrt{g}}{dg} \cdot \frac{dg}{dx}$

$$= \frac{1}{2\sqrt{g}} \cdot 2x$$

$$= \frac{1}{2\sqrt{x^2+y^2+\pi}} \cdot 2x = \frac{x}{\sqrt{x^2+y^2+\pi}}$$

Then differentiate

$$\frac{df}{dy} = \frac{d\sqrt{g}}{dg} \cdot \frac{dg}{dy}$$

$$= \frac{1}{2\sqrt{g}} \cdot 2y$$

$$= \frac{1}{2\sqrt{x^2+y^2+\pi}} \cdot 2y = \frac{y}{\sqrt{x^2+y^2+\pi}}$$

Therefore, the gradients (partial derivatives) of f with respect to x and y are:

$$= \frac{x}{\sqrt{x^2+y^2+\pi}} \text{ and } \frac{y}{\sqrt{x^2+y^2+\pi}}$$

1c.

Using the hint that a function with two variables f has a local maximum or minimum at (a, b) if $\partial f / \partial x(a, b) = 0$ and $\partial f / \partial y(a, b) = 0$),

$$\frac{df}{dx} \text{ and } \frac{df}{dy} \text{ must be } = 0$$

i.e

$$\frac{x}{\sqrt{x^2+y^2+\pi}} = 0, \text{ then } x = 0 \text{ since } \sqrt{x^2+y^2+\pi} \text{ will always be positive}$$

$$\frac{y}{\sqrt{x^2+y^2+\pi}} = 0 \text{ then } y = 0$$

Therefore $f(x,y)$ has a critical point at $(x,y)=(0,0)$

$$\text{i.e } f(0,0) = \sqrt{0^2 + 0^2 + \pi} = \sqrt{\pi}$$

We already know that $x^2 + y^2 \geq 0$

Adding π to both side,

$$= x^2 + y^2 + \pi \geq 0 + \pi$$

$$= x^2 + y^2 + \pi \geq \pi$$

Taking the square root

$$=\sqrt{x^2 + y^2 + \pi} \geq \sqrt{\pi}$$

Therefore, the minimal value of $f(x, y)$ is indeed $\sqrt{\pi}$.

2a.

$$\text{Since } X = \begin{pmatrix} 1 & x_1 \\ 1 & x_2 \\ \cdot & \cdot \\ \cdot & \cdot \\ 1 & x_n \end{pmatrix}, y = \begin{pmatrix} y_1 \\ y_2 \\ \cdot \\ \cdot \\ y_n \end{pmatrix} \text{ and } \omega = \begin{pmatrix} \omega_0 \\ \omega_1 \end{pmatrix}$$

the 1-norm of a vector $u = \|u\|_1 = \frac{1}{m} \sum_{i=1}^m |u_i|$

Using these 4, the cost function $L(x; \omega)$ rewritten using a matrix form is:

$$L(x; \omega) = \frac{1}{m} \|y - X\omega\|_1$$

2b.

Using the cost function:

$$L(x; \omega) = \frac{1}{m} \sum_{i=1}^m |y_i - (\omega_0 + \omega_1 x_i)|$$

Gradient with respect to ω_0 :

$$\text{Let } y_i - (\omega_0 + \omega_1 x_i) = a$$

$$L(x; \omega) = \frac{1}{m} \sum_{i=1}^m |a|$$

$$= \frac{d|a|}{d\omega_0} = \frac{a}{|a|} \cdot \frac{da}{d\omega_0}$$

= Since $a = y_i - (\omega_0 + \omega_1 x_i)$, we have

$$= \frac{da}{d\omega_0} = -1$$

$$\frac{d|a|}{d\omega_0} = \frac{a}{|a|} \cdot -1 = -\frac{a}{|a|}$$

$$\text{Therefore, } \frac{dL}{d\omega_0} = \frac{1}{m} \sum_{i=1}^m -\frac{a}{|a|}$$

$$= -\frac{1}{m} \sum_{i=1}^m (y_i - (\omega_0 + \omega_1 x_i))$$

Gradient with respect to ω_1 :

Let $y_i - (\omega_0 + \omega_1 x_i) = a$

$$L(x; \omega) = \frac{1}{m} \sum_{i=1}^m |a|$$

$$= \frac{d|a|}{d\omega_1} = \frac{a}{|a|} \cdot \frac{da}{d\omega_i}$$

= Since $a = y_i - (\omega_0 + \omega_1 x_i)$, we have

$$= \frac{da}{d\omega_i} = -x_i$$

$$\frac{d|a|}{d\omega_i} = \frac{a}{|a|} \cdot -x_i = -\frac{a}{|a|} x_i$$

$$\text{Therefore, } \frac{dL}{d\omega_i} = \frac{1}{m} \sum_{i=1}^m \frac{-a}{|a|} x_i$$

$$= -\frac{1}{m} \sum_{i=1}^m (y_i - (\omega_0 + \omega_1 x_i)) x_i$$

```
import numpy as np
import matplotlib.pyplot as plt
```

▼ Your Question3 here

We will see if we can find the results of question 1 using the gradient descent. a. Use the results of question 1.b) to find the couple (x, y) that minimizes $f(x, y)$, this time, with the gradient descent :

$$x_{n+1} = x_n - \mu \cdot \partial f / \partial x(x_n, y_n)$$

$$y_{n+1} = y_n - \mu \cdot \partial f / \partial y(x_n, y_n)$$

Use $\mu = 0.01$, $(x_0, y_0) = (7, 12)$ and define your number of iteration.

b. Plot the values of x and y during the iterative process.

c. Compute $f(x_{\text{final}}, y_{\text{final}})$ and compare with the result of question 1.

```
# Function f(x, y)
def f(x, y):
    return np.sqrt(x**2 + y**2 + np.pi)

# Gradients of f
def grad_f_x(x, y):
    return x / np.sqrt(x**2 + y**2 + np.pi)

def grad_f_y(x, y):
    return y / np.sqrt(x**2 + y**2 + np.pi)

# Gradient descent parameters
mu = 0.01
x0, y0 = 7, 12
iterations = 500
```

```
# Initialize variables
x, y = x0, y0
x_vals, y_vals = [x], [y]


# Gradient descent loop
for i in range(iterations):
    x_new = x - mu * grad_f_x(x, y)
    y_new = y - mu * grad_f_y(x, y)
    x, y = x_new, y_new
    x_vals.append(x)
    y_vals.append(y)
```

```
# Final values
x_final, y_final = x, y

# Plotting the values of x and y during the iterative process
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(x_vals, label='x values')
plt.xlabel('Iteration')
plt.ylabel('x value')
plt.legend()
plt.subplot(1, 2, 2)
plt.plot(y_vals, label='y values')
plt.xlabel('Iteration')
plt.ylabel('y value')
plt.legend()
plt.suptitle('Gradient Descent Process')
plt.show()
```



Gradient Descent Process



Generate

print hello world using rot13

Close

Generate is available for a limited time for unsubscribed users. [Upgrade to Colab Pro](#)

```
# Compute f(x_final, y_final)
f_final = f(x_final, y_final)
f_initial = f(x0, y0)

# Print the final results
print(f"Initial value of f(x0, y0): {f_initial}")
print(f"Final value of f(x_final, y_final): {f_final}")
print(f"Final (x, y): ({x_final}, {y_final})")
```

Initial value of f(x0, y0): 14.005055967527934
Final value of f(x_final, y_final): 9.128067105247087
Final (x, y): (4.5118271776279135, 7.73456087593357)


For Question4

```
data = np.loadtxt('Assignment-data_week3.csv', dtype=float, delimiter=',', converters={0:float , 1:float })
```

```
x = data[:,0]
y = data[:,1]
```

```
plt.scatter(x, y)
```

 <matplotlib.collections.PathCollection at 0x7a76cc2877f0>



```
# CONSTRUCT THE DESIGN MAT
X = np.vstack((np.ones_like(x), x)).T
```

```
def MAE(X, y, w_0, w_1):
    """
    Input
    X : design matrix
    y : the responses
    w_0, w_1 : initial guess

    Return
    loss (mean absolute error) : the loss evaluated at (t)

    """
    m = len(y)
    w = [w_0, w_1]
    loss = np.mean(np.abs(y - y_hat))

    return loss
```

```
def computeGradients(x, y, w_0, w_1):
    """
    Input
    x : the treatment or input
    y : the responses
    w_0, w_1 : initial guess

    Return
    grad_0, grad_1: gradients of the loss evaluated at (t)

    """
    # Define y_hat

    y_hat = w_0 + w_1 * x
    error = y - y_hat

    #IMPLEMENT THE gradients with respect to w_0 and w_1

    grad_w1 = -np.mean(np.sign(error))
    grad_w0 = -np.mean(np.sign(error) * x)
    return [grad_w0, grad_w1]
```

```
# initialize the weight
w_0 = 0
w_1 = 0

# Define the learning parameters
eta =0.01

num_iters =500

###
y_hat = w_0 + w_1 * x

lossVals = []
for i in range(num_iters):
    loss = MAE(X, y, w_0, w_1)
    lossVals.append(loss)

    # IMPLEMENT THE UPDATE using computeGradients(x, y, w_0, w_1)
    grad_w0, grad_w1 = computeGradients(x, y, w_0, w_1)

    w_0 -= eta * grad_w0
    w_1 -= eta * grad_w1

#Save the values of the loss in lossVals
print(f'Final w_0: {w_0}')
print(f'Final w_1: {w_1}')
```

```
Final w_0: 0.4895834146155935
Final w_1: 0.01
Final w_0: 0.979166829231187
Final w_1: 0.02
Final w_0: 1.4687502438467805
Final w_1: 0.03
Final w_0: 1.958333658462374
Final w_1: 0.04
Final w_0: 2.4479170730779676
Final w_1: 0.05
Final w_0: 2.937500487693561
Final w_1: 0.060000000000000005
Final w_0: 3.4270839023091546
Final w_1: 0.07
Final w_0: 3.916667316924748
Final w_1: 0.08
Final w_0: 4.406250731540341
Final w_1: 0.09
Final w_0: 4.895834146155934
Final w_1: 0.09999999999999999
Final w_0: 5.385417560771527
Final w_1: 0.10999999999999999
Final w_0: 5.87500097538712
Final w_1: 0.11999999999999998
Final w_0: 6.3645843900027135
Final w_1: 0.12999999999999998
Final w_0: 6.8541678046183065
Final w_1: 0.13999999999999999
Final w_0: 7.3437512192339
Final w_1: 0.15
Final w_0: 7.833334633849493
Final w_1: 0.16
Final w_0: 8.322918048465086
Final w_1: 0.17
Final w_0: 8.812501463080679
Final w_1: 0.18000000000000002
Final w_0: 9.302084877696272
Final w_1: 0.19000000000000003
Final w_0: 9.791668292311865
Final w_1: 0.20000000000000004
Final w_0: 10.281251706927458
Final w_1: 0.21000000000000005
Final w_0: 10.770835121543051
Final w_1: 0.22000000000000006
Final w_0: 11.260418536158644
Final w_1: 0.23000000000000007
Final w_0: 11.750001950774237
Final w_1: 0.24000000000000007
Final w_0: 12.23958536538983
Final w_1: 0.25000000000000006
Final w_0: 12.729168780005423
```



```
Final w_1: 0.26000000000000006
Final w_0: 13.218752194621016
Final w_1: 0.2700000000000001
Final w_0: 13.70833560923661
Final w_1: 0.2800000000000001
Final w_0: 14.197919023852203
```


```
#####
# visualizing the loss function
#####

plt.plot(lossVals)
plt.xlabel('Iteration')
plt.ylabel('Mean Absolute Error')
plt.title('Error Evolution During Gradient Descent')
plt.show()

# Plot the regression line
plt.scatter(x, y, label='Data')
plt.plot(x, w_0 + w_1 * x, color='red', label='Fitted Line')
plt.xlabel('x')
plt.ylabel('y')
plt.title('Linear Regression Fit')
plt.legend()
plt.show()
```



Error Evolution During Gradient Descent



Linear Regression Fit

